

Tower Dropper

Group Z

Paolo Deidda (paolo.deidda@usi.ch)
Francesco Costa (francesco.costa@usi.ch)
Raffaele Perri (raffaele.perri@usi.ch)

<https://github.com/USI-Projects-Collection/FinalProject-Robotics.git>

Contents

1	Introduction	1
1.1	Model and Scene	1
2	Workflow of the Project	1
2.1	Global Path Planning	1
2.1.1	Objectives and Requirements	1
2.1.2	Algorithm Selection and Design	1
2.1.3	Map Representation, Visualization & Preprocessing	2
2.1.4	ROS2 Integration	2
2.1.5	Issues Encountered & Troubleshooting	2
2.2	Tower Detection and Interaction System	2
2.2.1	System Overview	2
2.2.2	Sensor Integration and Data Processing	3
2.2.3	State Machine Implementation and Control Logic	3
2.3	Weak-Spot Detection & Shooting	5
2.3.1	Tower Alignment and Targeting System	5
2.3.2	Alignment Detection	5
2.3.3	Projectile Launch System	5
3	System Integration	6
3.1	ROS 2 Node Graph & Communication Channels	6
3.2	Integration Strategy and Version Control	6
3.3	Observed Failure Modes & Mitigations	6
4	Results	6
5	Conclusion & Future Improvements	6
5.1	Potential Improvements	6
5.2	Unexplored Ideas	7

Setup

To run the code follow the setup and execution instructions provided in the `README.md` file included in this repository.

1 Introduction

The goal of this project was to control a robot with the objective of knocking towers to the ground through the action of a mounted turret. A successful implementation of such a system requires a number of different individual task, including:

- Exploration of the environment with the goal of finding the towers
- Detection of the tower and of its weak-spot
- Alignment with the tower
- Simulation of the shooting process

All of these different problems required diverse solutions and brought distinct difficulties, making the project challenging but all the more rewarding.

1.1 Model and Scene

The project use a RobomasterS1 robot model, with the addition of four proximity sensors, placed respectively (Figure 1):

- **Range_0 (Center-right):** Primary sensor for orbit distance maintenance during left-side orbital maneuvers. Positioned at the robot's right quadrant with a detection range up to 10 meters.
- **Range_1 (Front-right):** Primary sensor for initial tower detection and forward-right obstacle detection. Critical for the search phase and collision avoidance.
- **Range_2 (Center-left):** Secondary sensor utilized during obstacle avoidance protocols, enabling the robot to maintain safe distances from obstacles while navigating around them.
- **Range_3 (Front-left):** Forward collision prevention sensor that triggers emergency avoidance behaviors when obstacles are detected within critical thresholds.



Figure 1: RobomasterS1 robot model with camera and proximity sensors.

The robot is equipped with a camera (`/rm0/camera/image_color`) used for visual processing tasks, such as weak-spot detection and tower

identification.

The scene is constructed with a 3D model of the RobomasterS1 robot placed in the center of a flat terrain. The environment includes four towers, each colored in red to indicate that they are active and can be targeted. The towers are positioned at the corners of the scene.




Figure 2: Scene depicting environment.

Additionally there is a white pillar resembling a tower which is used to showcase the obstacle avoidance capabilities of the robot.

2 Workflow of the Project

2.1 Global Path Planning

2.1.1 Objectives and Requirements

The goal of our global planner is to compute, at runtime, a collision-free route from the robot's current pose to the next tower in the Tower Dropper scenario. The planner must respect:

- **Static environment:** a $5\text{ m} \times 5\text{ m}$ occupancy grid (5 cm resolution) generated by our mock map publisher.
- **Robot footprint:** approximated as a circle of radius 0.3 m, enforced via obstacle inflation.
- **Real-time constraints:** paths should be available within a few hundred milliseconds to allow seamless integration with our local controller and to react promptly when the mock publisher advances to the next tower.

2.1.2 Algorithm Selection and Design

We surveyed classic grid-based planners (Dijkstra for guaranteed optimality, RRT for sampling-based routes) but settled on **A*** on an occupancy grid for its simplicity and determinism. In our `PathPlannerNode`:

1. **Grid representation:** a 2D NumPy array of occupancy values ($0 = \text{free}$, $\geq 50 = \text{occupied}$).
2. **Search:** an 8-connected A* using a min-heap (`heapq`) keyed by $f = g + h$, where g is accumulated path cost and h the Euclidean distance to goal. Orthogonal

moves cost 1; diagonal moves cost $\sqrt{2}$. Corner-cutting checks prevent invalid diagonals.

3. **Path pruning:** after finding the raw A* path, we remove unnecessary waypoints by testing straight-line visibility with a Bresenham-based `_has_los` check.

2.1.3 Map Representation, Visualization & Pre-processing

To debug and tune our planner, we relied heavily on RViz2, subscribing to the `/map` and `/plan` topics to inspect both the occupancy grid and computed paths in real time. Figures 3 and 4 illustrate this process.




Figure 3: Initial pathfinding: left—raw A* solution; right—after line-of-sight pruning.

Subsequently, we introduced obstacle inflation to enforce the robot’s safety margin and avoid corner collisions:




Figure 4: Pathfinding on the inflated map (beige overlay): the planner keeps clear of obstacle edges.

Preprocessing steps:

- Reshape `OccupancyGrid.data` into a 2D array.
- Inflate each occupied cell by `inflate_radius = int(0.3 / resolution)` to mark neighboring cells as occupied, ensuring collision clearance.

2.1.4 ROS2 Integration

The planner is implemented as a ROS 2 node with:

Subscriptions: `/map` (`OccupancyGrid`), `/goal_pose` (`PoseStamped`), `/rm0/odom` (`Odometry`), `/go_again`, `/goal_reached` (`Bool`).

Publishers: `/plan` (`nav_msgs/Path`), `/current_pose` (`PoseStamped`), `/rm0/cmd_vel` (`Twist`), `/goal_reached` (`Bool`).

TF Broadcast: emits an `odom → base_link` transform each update, plus a `static_transform_publisher` for `map → odom` to align frames in RViz.

Namespace handling: launch with `--namespace <robot>` so one codebase runs on multiple robots unchanged.

2.1.5 Issues Encountered & Troubleshooting

TF frame misalignment *Symptom:* paths appeared shifted in RViz. *Diagnosis:* no `map → odom` static transform in `/tf`. *Solution:* added a `static_transform_publisher` in the launch file to broadcast an identity `map → odom`.

Unknown-cell “blockades” *Symptom:* planner routed through unexplored (−1) cells on SLAM maps. *Diagnosis:* we treated −1 as free in our occupancy threshold. *Solution:* in `map_cb`, mark any `grid[y][x] < 0` as 100 before inflation, treating unknown as occupied.

Occupied goal cell handling *Symptom:* when running the tower-knockdown process, the robot sometimes started inside a region forbidden by the inflated map, causing “no path found.” *Diagnosis:* both start and goal cells were disallowed by inflation. *Solution:* implemented `nearest_free(goal_cell)`, which performs an orthogonal BFS to locate the closest free cell; we then replace the goal with that cell and re-plan automatically.

2.2 Tower Detection and Interaction System

2.2.1 System Overview

The tower detection and interaction system represents a sophisticated autonomous targeting solution. The system integrates multiple sensor modalities including range sensors, RGB camera, and odometry to enable a mobile robot to autonomously locate, approach, orbit, and engage tower targets in a dynamic environment. The core architecture follows a finite state machine approach with six primary operational states: *find_tower*, *position_for_orbit*, *orbit_tower*, *avoid_obstacle*, *align_tower*, and *shoot_tower*. This state-based design ensures robust behavior transitions and clear operational logic flow through well-defined transition criteria and sensor-based decision making. The main control loop executes at 60 Hz, providing real-time responsiveness while maintaining computational efficiency.

Key ROS2 topics include velocity commands (`cmd_vel`), odometry data (`odom`), individual range sensor readings (`/rm0/range_0-3`), and camera imagery (`/rm0/camera/image_color`).

2.2.2 Sensor Integration and Data Processing

Range Sensor Configuration and Calibration

The system utilizes four range sensors positioned strategically around the robot platform to provide environmental awareness (as shown in Figure 1).

Each sensor operates through dedicated ROS2 subscribers that continuously update distance measurements at the sensor's native frequency. The system implements sensor fusion by maintaining a real-time comparison between multiple sensor readings to determine the most reliable distance measurements and detect sensor failures or environmental interference. Readings exceeding 10 meters are considered invalid and replaced with the maximum range value. The system also implements a smoothing algorithm to reduce measurement noise:

$$d_{filtered}(t) = \alpha \cdot d_{raw}(t) + (1 - \alpha) \cdot d_{filtered}(t - 1) \quad (1)$$

where $\alpha = 0.8$ provides a balance between responsiveness and stability.

Computer Vision System Architecture

The vision system employs a comprehensive computer vision pipeline optimized for real-time tower detection and tracking. The process begins with image acquisition from the robot's RGB camera, followed by color space conversion and multi-stage filtering.

Color Space Processing: The system converts incoming BGR images to HSV color space to achieve robust color-based detection under varying lighting conditions. HSV provides superior color constancy compared to RGB, particularly important for outdoor or variable lighting scenarios.

Multi-Range Color Filtering: To account for the bimodal distribution of red pixels in HSV space, the system implements dual-range thresholding:

$$\text{Mask}_1 = \text{inRange}(\text{HSV}, [0, 100, 100], [10, 255, 255]) \quad (2)$$

$$\text{Mask}_2 = \text{inRange}(\text{HSV}, [160, 100, 100], [179, 255, 255]) \quad (3)$$

$$\text{Mask}_{final} = \text{Mask}_1 \cup \text{Mask}_2 \quad (4)$$

This approach captures both low-end ($0 - 10^\circ$) and high-end ($160 - 179^\circ$) red hues while maintaining robust detection across different illumination conditions.



Figure 5: Multi-range color filtering process for red tower detection.

Tower Position Calculation: The system calculates tower position using spatial moments of the binary mask:

$$M_{pq} = \sum_{x,y} x^p y^q \cdot \text{Mask}(x, y) \quad C_x = \frac{M_{10}}{M_{00}} \quad (5)$$

$$p_{tower} = \frac{C_x - W/2}{W/2} \quad (6)$$

where M_{pq} represents the (p, q) -th moment, C_x is the centroid x-coordinate, W is the image width, and $p_{tower} \in [-1, 1]$ represents the normalized tower position.

Tower Visibility Assessment: The system implements a visibility metric based on the ratio of red pixels to total image pixels:

$$R_{red} = \frac{\sum_{x,y} \text{Mask}_{final}(x, y)}{W \times H} \quad (7)$$

Tower visibility is confirmed when $R_{red} > \tau_{visibility}$, where the threshold $\tau_{visibility}$ is dynamically adjusted based on environmental conditions and historical detection performance.

2.2.3 State Machine Implementation and Control Logic

State Transition Architecture

The finite state machine implementation employs a hierarchical structure with clearly defined transition conditions and state-specific behaviors (Figure 6). Each state maintains internal variables for tracking progress and implementing timeouts to prevent infinite loops or deadlock conditions. State transitions are triggered by combinations of sensor readings, timing constraints, and task completion flags. The system implements hysteresis in transition conditions to prevent oscillation between states due to sensor noise or borderline threshold conditions.




Figure 6: Robot state machine diagram illustrating the tower detection and interaction workflow.

Tower Discovery Phase: *find_tower* State

The robot executes counterclockwise rotation at $\omega_{search} = 0.5$ rad/s while continuously monitoring the front-right range sensor (Range.1). This angular velocity provides optimal balance between search speed and sensor update rates, ensuring no potential targets are missed during rotation. **Detection Criteria:** Tower detection occurs when the front-right sensor measurement satisfies:

$$d_{Range1} \leq d_{target} + \delta_{tolerance} \quad (8)$$

where $d_{target} = 2.0$ meters represents the desired orbital radius and $\delta_{tolerance} = 0.1$ meters provides measurement uncertainty accommodation.

Multi-Sensor Validation: Upon initial detection, the system performs cross-validation using adjacent sensors to confirm target authenticity and eliminate false positives from environmental artifacts or sensor noise. The validation process requires consistent readings across multiple sensor cycles to prevent premature state transitions.

Orbital Positioning Phase: *position_for_orbit* State

This critical state ensures proper spatial relationship between robot and tower before initiating orbital maneuvers.

Geometric Positioning Strategy: The system executes a controlled rotation to position the detected tower relative to the robot's right side, enabling subsequent left-side orbital motion. The positioning maneuver uses angular velocity $\omega_{position} = 0.5$ rad/s while monitoring the back-right sensor (Range.0).

Position Validation Criteria: Successful positioning requires the back-right sensor to detect the tower within acceptable bounds:

$$d_{Range0} \leq d_{target} + \delta_{position} \quad (9)$$

where $\delta_{position} = 0.5$ meters provides sufficient tolerance for initial orbital entry while maintaining proximity to the target distance.

Timeout and Recovery Mechanisms: The state implements timeout protection to prevent infinite positioning attempts. If positioning cannot be achieved within a pre-determined time window, the system reverts to the search state with modified search parameters.

Orbital Navigation Control: *orbit_tower* State

The orbital control system represents the most sophisticated aspect of the navigation architecture, combining multiple control loops for stable circular motion around the detected tower.

Multi-Loop Control Architecture: The orbital controller implements two primary control loops operating in parallel:

1. **Distance Regulation Loop:** Maintains constant radial distance using range sensor feedback
2. **Angular Tracking Loop:** Centers tower in camera field of view using vision feedback

Distance Control Implementation: The distance regulation employs PID control with anti-windup protection:

$$e_d(t) = d_{measured} - d_{target} \quad (10)$$

$$u_d(t) = K_p e_d(t) + K_i \int_{t_0}^t e_d(\tau) d\tau + K_d \frac{de_d(t)}{dt} \quad (11)$$

The integral term includes saturation limits to prevent windup:

$$\int_{t_0}^t e_d(\tau) d\tau = \begin{cases} I_{\max} & \text{if } \int e_d(\tau) d\tau > I_{\max} \\ I_{\min} & \text{if } \int e_d(\tau) d\tau < I_{\min} \\ \int e_d(\tau) d\tau & \text{otherwise} \end{cases} \quad (12)$$

where $I_{\max} = 5.0$ and $I_{\min} = -5.0$ prevent excessive integral buildup.

Velocity Coordination: The orbital motion combines base velocities with correction terms from both control loops:

$$\begin{aligned} v_{linear} &= v_{base} \cdot f_{distance}(e_d) \omega_{angular} \\ &= \omega_{base} + \alpha \cdot u_d + \beta \cdot p_{tower} \end{aligned} \quad (13)$$

where:

$$f_{distance}(e_d) = \begin{cases} 0.8 & \text{if } e_d > 0.2 \text{ (too far)} \\ 1.2 & \text{if } e_d < -0.2 \text{ (too close)} \\ 1.0 & \text{otherwise} \end{cases} \quad (14)$$

The velocity adaptation function $f_{distance}$ provides reactive speed adjustments based on distance error magnitude, enhancing convergence to the target orbital radius. The system implements comprehensive velocity limiting to ensure safe operation: $v_{linear} \in [0.05, 0.3]$ m/s $\omega_{angular} \in [-\infty, -0.1]$ rad/s. The angular velocity constraint ensures consistent counterclockwise motion while preventing excessive rotation rates that could destabilize the orbital behavior.

Obstacle Avoidance Protocol: *avoid_obstacle* State

The obstacle avoidance system implements a dual-objective navigation strategy that maintains safe distances from obstacles while preserving awareness of

the primary tower target.

Obstacle Detection and Classification: Obstacle detection triggers when front sensors detect objects within the critical threshold.

Secondary Orbital Behavior: Upon obstacle detection, the system initiates a secondary orbital pattern around the detected obstacle using the back-left sensor (Range_2) for guidance.




Figure 7: Robot avoiding the top obstacle encountered during orbital navigation.

Return Navigation Logic: The system implements logic for returning to the primary tower after obstacle avoidance. The return condition requires:

$$(d_{Range2} > d_{clearance}) \wedge (d_{Range0} < d_{tower,max}) \wedge \text{flag}_{almost_avoided} \quad (15)$$

where $d_{clearance} = 10.0$ meters indicates obstacle clearance and $\text{flag}_{almost_avoided}$ prevents premature return attempts.

2.3 Weak-Spot Detection & Shooting

2.3.1 Tower Alignment and Targeting System

The system requires alignment algorithm to be used in order to detect the weak side of the tower which, when working with rectangular bases, simply equates to the largest side. The problem could easily be solved by coloring different sides of the tower or by including a special mark on the target side. This solution would harm the implementation on the long run, as we would always be forced to explicitly distinguish sides by altering the visual properties of towers, leading to more constrained and less realistic results.

2.3.2 Alignment Detection

The method we developed to detect the weak side of the tower makes use of the binary mask already in use to assess the position of the tower. In this case we take advantage

of the discretization limits of pixels. The easiest, and surprisingly reliable, method to detect whether the robot is looking at a flat face of the tower relies in fact on the distribution of colored pixels in the mask, in particular around the base of the tower. If the number of continuous colored pixels in the first row of the mask containing any non-black pixel (starting from the bottom) is around as high as the highest number of continuous pixels in a row of the mask, we can confidently say we are looking at a flat face of the tower. Visually this would mean that in the binary mask we see a sharp edge at the bottom of the tower, instead of a 'pixelated' line.

Alternative Alignment Detection

In the search for the best alignment detection algorithm we tried a number of different options, one of which we kept as an alternative. This particular solution employs the `findContours` function of openCV which, given a mask, returns the contours of an image. With this information we are then able to detect situations in which the contour is a perfect rectangle and discriminate between small and large faces by looking at the ratio between height and base of the contour. This solution proved harder to debug since it makes use of functions of openCV, giving us less freedom in the implementation. For this reason we decided to stick with the previously described solution.

Sensor-Based Geometric Validation

Once the robot is confident of being parallel to the largest side of the tower we follow an iterative process in order to turn the robot until it looks straight towards its target. The alignment process follows a systematic approach:

1. **Initial Rotation:** Controlled rotation at $\omega_{align} = -0.3$ rad/s to scan for geometric variations
2. **Asymmetry Detection:** Continuous monitoring of range sensor differences to identify tower edges
3. **Symmetry Convergence:** Fine positioning adjustments to achieve bilateral sensor equality
4. **Alignment Correction:** Final correction of the alignment following rotation direction

2.3.3 Projectile Launch System

The *shoot_tower* state executes the engagement sequence through integration with the CoppeliaSim physics engine. The system calculates projectile spawn position using coordinate transformation:

$$\mathbf{R} = \begin{bmatrix} \cos(\psi) & -\sin(\psi) & 0 \\ \sin(\psi) & \cos(\psi) & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (16)$$

$$\mathbf{p}_{projectile} = \mathbf{p}_{robot} + \mathbf{R} \cdot \mathbf{t}_{blaster} \quad (17)$$

where ψ represents the robot's roll angle and $\mathbf{t}_{blaster} = [0.2, 0, 0.2]^T$ meters defines the blaster offset from the

robot center. The projectile receives initial velocity $\mathbf{v}_{projectile} = \mathbf{R} \cdot [3.0, 0.0, 3.0]^T$ m/s, providing forward momentum with upward trajectory compensation for ballistic flight.

3 System Integration

3.1 ROS 2 Node Graph & Communication Channels

The communication between the three ROS2 nodes in the system — `MockMapPublisher`, `PathPlannerNode`, and `ControllerNode` — is established through various ROS topics. Here’s an overview of their coordination logic:

At initialization, the `MockMapPublisher` and `PathPlannerNode` nodes sets the internal flag `self.go_again = True` and the `ControllerNode` sets `self.goal_reached = False`. These initial states allow both the `MockMapPublisher` and `PathPlannerNode` to begin processing immediately, while the `ControllerNode` waits for confirmation that a goal has been reached. During operation, once the `PathPlannerNode` reaches its goal, it publishes a velocity command `cmd_vel = (0.0, 0.0)` to bring the robot to a halt. Then publishes the following message to notify other nodes of goal completion:

```
self.pub_goal_reached.publish(Bool(data=True))
```

and sets `self.go_again = False` to indicate that no further planning is needed until explicitly requested again. The `ControllerNode`, which subscribes to the `/goal_reached` topic, receives this message and updates its internal flag (`self.goal_reached = True`). Upon completing its own task (such as analyzing sensor data or interacting with the tower), it publishes:

```
self.turn_ended.publish(Bool(data=True))
```

on the topic `/go_again`. Both the `MockMapPublisher` and `PathPlannerNode`, which subscribe to `/go_again`, receive this signal and set `self.go_again = True`. This loop allows the system to autonomously progress through a series of goals or states in a coordinated manner. This design ensures that the map publisher and planner are event-driven and coordinated with the controller logic, enabling robust and synchronized robot behavior across discrete tasks.

3.2 Integration Strategy and Version Control

3.3 Observed Failure Modes & Mitigations

4 Results

The final implementation of the project is able to produce results matching our initial set of requirements in a satisfactory way. The robot in fact capable of, thorough path

planning, reach the position of all towers on the scene. Once the robot is in the vicinity of the tower it then proceeds to orbit around the perimeter of the target in search of the weak-spot (largest flat side). Finally the robot aligns itself to the face of the tower and the action of shooting a projectile towards it is simulated by creating a sphere with adequate mass and velocity.




Figure 8: Snapshot of the robot after shooting a projectile.

A complete execution demo can be found at this link.

5 Conclusion & Future Improvements

5.1 Potential Improvements

In conclusion, while the project successfully implements all goal behaviors, there are a number of possible improvements which could be done to it:

- First and foremost additional work could be done to improve robustness of the system. In its current state the control of the robot is in fact limited to certain parameters, such as the static definition of the scene (position of the towers and obstacles). Additionally, both avoidance of obstacles and weak-spot detection don’t take into consideration certain edge cases we decided to purposely avoid (obstacles too close to the tower and situations in which the robot orbits too close to it).
- Regarding the exploration of the environment we make use of path planning to find the position of all target towers without the use of sensors (visual or distance). A more advanced solution could employ the use of said sensors and dynamically map the environment around the robot, making it possible to generate almost randomly the scenes.
- Finally, with the current implementation we are not making use of the moveable gimbal supporting the turret. It could be possible through its functionalities to improve the detection and tracking of the towers, leading to a more natural and satisfying result. This would completely remove the need to realign the robot to the tower after the flat side is detected.

5.2 Unexplored Ideas

Given the very general description of the problem there are also a number of possible unexplored ideas which could fuel future projects. One such idea would be to include multiple robots in the scene (instead of towers) and have them target each other in an all out brawl. The physics engine provided with CoppeliaSim would allow for an extremely satisfying simulation. It could also be possible to think of a situation where the robot is not moving (or barely) and the focus is instead on the control of the mounted turret, which would need to follow moving targets and predict their trajectory before shooting projectiles at them.